

# Your Claude Model Playbook

Which model (and effort level) to reach for — mapped to how you actually work

PRO PLAN (~\$20/MO) CODING WRITING & EDITING EVERYDAY CHAT AUG 2026

**The 30-second version.** Make **Sonnet 5** your daily driver — it handles most of your chatting, writing, and routine coding while barely touching your usage allowance. Step up to **Opus** (5 or 4.8) only when a coding or reasoning task is genuinely hard. Save **Fable 5** for your rare toughest problems — on Pro it draws from metered usage credits, so it's the one that can actually run you dry. Drop to **Haiku 4.5** for quick throwaway questions. And treat **Effort** as a second dial: higher effort = better on hard problems, but slower and more of your allowance spent.

## The line-up at a glance

| MODEL              | TIER                   | BEST FOR                                                         | RELATIVE COST*       |
|--------------------|------------------------|------------------------------------------------------------------|----------------------|
| Fable 5            | Top (Mythos-class)     | Your toughest, highest-stakes problems & creative writing        | Highest · ~10× Haiku |
| Opus 5             | Flagship               | Hard agentic coding, deep reasoning, complex builds              | High · ~5× Haiku     |
| Opus 4.8           | Flagship (app default) | Same lane as Opus 5; the model your picker currently defaults to | High · ~5× Haiku     |
| Sonnet 5           | Balanced               | Everyday work — most coding, writing & editing. Your default.    | Low · ~2× Haiku      |
| Haiku 4.5          | Fast                   | Quick answers, simple lookups, short rewrites                    | Lowest · baseline    |
| Opus 4.7 / 4.6 / 3 | Legacy                 | Older Opus versions — no reason to pick over Opus 5/4.8          | High (skip)          |
| Sonnet 4.6         | Legacy                 | Previous Sonnet — Sonnet 5 is better and cheaper                 | Medium (skip)        |

\*You're on Pro, so you don't pay per token — you pay a flat subscription and draw from a usage allowance. "Relative cost" is the per-token API price used purely as a proxy for how fast each model eats that allowance. Higher multiple = you hit limits sooner.

## The models, one by one

### Sonnet 5 Your default driver

API proxy: \$2 / \$10 per M tokens

The "most efficient for everyday tasks" option — and the current default on Free/Pro for good reason. Fast, strong at code and prose, and light on your allowance so you can run it all day.

- Reach for it first for: routine coding, drafting & editing, summaries, everyday questions.
- Only escalate when it visibly struggles — tangled bugs, deep architecture, or unusually nuanced writing.

### Haiku 4.5 The quick one

API proxy: \$1 / \$5 per M tokens

Fastest and cheapest. "Fastest for quick answers." Perfect when you want a snappy reply and don't need heavy reasoning.

- Great for: quick facts, syntax reminders, short rewrites, "what does this error mean," brainstorming.
- Don't use it for: multi-step logic, long documents, or anything where quality really matters.

The flagship reasoning/coding tier. Opus 5 is the newer one; Opus 4.8 is what your app currently defaults to — they sit in the same lane, so either is a fine "hard problems" button. Noticeably more capable than Sonnet on genuinely difficult work, but heavier on your allowance.

- Reach for it when: a bug has beaten Sonnet, you're designing a system, refactoring across many files, or reasoning through something intricate.
- Pair with **High** or **Extra** effort for the hardest cases.

### Fable 5 The top tier — use sparingly

API proxy: \$10 / \$50 per M tokens · "Requires usage credits" on Pro

Anthropic's most capable model, "for your toughest challenges," and the strongest at demanding creative writing. The catch for you: on Pro it's **metered by usage credits** (it's included outright only on Max). That badge in your picker is the warning label — this is the model that can actually cost you extra or run out.

- Worth it for: a genuinely gnarly problem Opus couldn't crack, or a high-stakes piece of writing where tone and craft really matter.
- Not worth it for: everyday coding, chatting, or routine drafts — you'd burn credits for little gain.

## The Effort dial (Low → Max → Thinking)

Effort controls how much the model reasons before answering. It's independent of which model you pick, and it moves cost as much as the model choice does — higher effort means more thinking, better answers on hard problems, but slower replies and more of your allowance spent.

| SETTING                         | WHEN TO USE IT                                                                                             |
|---------------------------------|------------------------------------------------------------------------------------------------------------|
| <b>Low</b>                      | Chat, quick lookups, simple edits. Fast and cheap — don't overthink easy things.                           |
| <b>Medium</b>                   | Solid everyday default for most writing and routine coding.                                                |
| <b>High</b> (the app's default) | Real problem-solving — complex code, careful arguments, layered reasoning.                                 |
| <b>Extra / Max</b>              | The hardest debugging, architecture, or analysis. Slowest and most costly — reserve them.                  |
| <b>Thinking</b>                 | Lets it think longer still. Best for the occasional problem that genuinely needs deep, extended reasoning. |

Rule of thumb: match effort to difficulty, not to importance. A simple email to an important person still only needs Low–Medium.

## How this maps to *your* three kinds of work

### CODING & BUILDING APPS

**Default:** Sonnet 5 at Medium/High effort — handles most features, edits, boilerplate, and "explain this code."

**Escalate to:** Opus 5 or Opus 4.8 (High → Extra) when a bug won't die, you're refactoring across files, or designing architecture.

**Top gear:** Fable 5 only for a truly stubborn algorithmic/architectural problem Opus can't solve — then switch back (credits).

**Quick stuff:** Haiku 4.5 for syntax questions, quick snippets, or decoding an error message.

Tip: start a hard debugging session on Opus + High rather than grinding many Sonnet attempts — one strong pass often costs less of your allowance than ten weak ones.

### WRITING & EDITING

**Default:** Sonnet 5 at Low/Medium — drafting, editing, tightening, summarizing, tone tweaks. The sweet spot for most prose.

**Escalate to:** Opus (High) for complex, structured, or argument-heavy pieces — reports, proposals, anything with a lot of moving parts.

**Top gear:** Fable 5 for high-stakes creative writing where craft and voice really matter — it's literally the model built for the toughest, most creative challenges. Reserve it (credits).

**Quick stuff:** Haiku 4.5 for subject lines, quick rephrases, grammar fixes.

EVERYDAY Q&A / CHAT

**Default:** Haiku 4.5 or Sonnet 5 at Low effort — fast, plenty smart, and barely touches your allowance.

**Avoid:** Opus and Fable here. Using a flagship for casual questions is the single easiest way to waste your usage — you pay heavy-model cost for no real benefit.

If a "simple" question turns into a real problem mid-conversation, that's your cue to switch up a model or raise effort — not before.

Cost, the Pro-plan way to think about it

On Pro there's no per-message price tag. You pay your flat ~\$20/month and draw from a usage allowance that refills on a rolling basis. So "cost" really means **how fast a choice uses that allowance** — and two things drive it: the **model** and the **effort level**. Fable 5 is the special case: it's separately **metered by usage credits** on Pro, so it's the one that can run out or cost extra.

| CHOICE                         | WHAT IT DOES TO YOUR ALLOWANCE                                                        |
|--------------------------------|---------------------------------------------------------------------------------------|
| Haiku 4.5, Low effort          | Barely a dent — chat freely.                                                          |
| Sonnet 5, Medium               | Light — your comfortable all-day default.                                             |
| Opus 5 / 4.8, High             | Several times heavier — use deliberately for hard work.                               |
| Fable 5                        | Heaviest, and drawn from metered credits on Pro — reserve for the rare toughest task. |
| Any model + Extra/Max/Thinking | Multiplies the above — the effort dial matters as much as the model.                  |

Quick decision cheat-sheet

|                                                         |                                                     |
|---------------------------------------------------------|-----------------------------------------------------|
| Just chatting or a quick question?                      | Haiku 4.5 or Sonnet 5 · Low                         |
| Normal coding, writing, or editing?                     | Sonnet 5 · Medium                                   |
| Hard bug, big refactor, or complex reasoning?           | Opus 5 or 4.8 · High → Extra                        |
| Toughest problem Opus can't crack / high-craft writing? | Fable 5 (watch credits) · High+                     |
| Seeing "Requires usage credits"?                        | That's Fable — fine occasionally, not for daily use |

**A note on freshness & sources.** Anthropic ships models fast and this lineup shifted a lot in mid-2026, so treat specifics as current-as-of-August-2026, not permanent. Your app's default is Opus 4.8; Opus 5 and Sonnet 5 are the newer siblings. Fable-on-Pro being metered by usage credits (and included on Max) is Anthropic's stated plan setup. Per-token figures are the public API rates, shown only to compare relative cost — they are not what you pay on Pro. For the definitive, live details on your plan's limits and Fable access, see [support.claude.com](#); for model pricing, [docs.claude.com](#). When in doubt, the in-app picker's own descriptions ("for your toughest challenges," "most efficient for everyday tasks," "fastest for quick answers") are a reliable quick guide.